

K-NN

k Nearest Neighbour Classifier

Eager and Lazy Learners

Lazy learners: Save all data from training, use it for classifying



(The learner was lazy, classifier had to do the work)

Eager learners: Build a (compact) model/structure during training, use the model for classification.

(The learner was eager/worked harder, classifier had simple life)

kNN is a lazy learner.

k-NN: Basic Idea



<https://towardsdatascience.com/knn-using-scikit-learn-c6bed765be75>

Distance Metrics

- Minkowski distance (L norm)

$$d_p(\mathbf{x}, \mathbf{y}) = \sqrt[p]{\sum_{i=1}^m |x_i - y_i|^p}$$

- Most popular: Euclidean Distance L2 (p=2)

$$d_E(\mathbf{x}, \mathbf{y}) = \sqrt{\sum_{i=1}^m (x_i - y_i)^2}$$

- e.g., $d((1,2),(3,4))=2.83$
- Manhattan distance L1 (p=1) $d_M(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^m |x_i - y_i|$
 - e.g., $d((1,2),(3,4))=4$

k-NN Algorithm

Training algorithm:

- For each training example $\langle x_j, f(x_j) \rangle$, where $f(x_j)$ is class of x_j , add the example to the list *training_examples*.

Classification algorithm:

- Given a query instance x_q to be classified, Let (x_1, \dots, x_k) be the k instances from *training_examples* that are *nearest* to x_q by the distance function. Let V be the set of classes for the k training examples.
- Return $class(x_q) = \arg \max_{v \in V} \sum_{i=1}^k \delta(v, f(x_i))$

Where $\delta(a,b) = 1$ if $a = b$ and 0 otherwise.

k-NN Example

outlook	temperature	humidity	windy	play	distance
sunny	hot	high	FALSE	no	2 (*)
sunny	hot	high	TRUE	no	1 (*)
overcast	hot	high	FALSE	yes	3
rainy	mild	high	FALSE	yes	3
rainy	cool	normal	FALSE	yes	3
rainy	cool	normal	TRUE	no	2 (*)
overcast	cool	normal	TRUE	yes	2 (*)
sunny	mild	high	FALSE	no	2 (*)
sunny	cool	normal	FALSE	yes	2 (*)
rainy	mild	normal	FALSE	yes	4
sunny	mild	normal	TRUE	yes	2 (*)
overcast	mild	high	TRUE	yes	2 (*)
overcast	hot	normal	FALSE	yes	4
rainy	mild	high	TRUE	no	1 (*)

We have Play Tennis data on the left

And we have new instance: $\langle Outlk = sun, Temp = cool, Humid = high, Wind = true \rangle$

- In order to calculate distance, each feature is coded as 1 or 0, so feature space has 10 dimensions. So the $x_i - x_j$ is 0 if $x_i = x_j$, or 1 otherwise.
- By Manhattan distance we need to find 5 nearest neighbours.
- But here we have lots of ties, therefore we can choose randomly to make 5, or count all ties (*).
- So we have 5 “no” and 4 “yes”, by voting, decision is NO

(but notice that in this dataset, features can also be coded as ordinal, e.g., outlook(sunny, overcast, rainy)->(3,2,1). The problem is: what number do we assign to those features?)

How to choose k for k -NN

- $k=1$ perfectly separates training data, so low bias but high variance
- By increasing the number of neighbours k we increase bias and decrease variance (what happens when $k = m$? m is the number of observations)
- Ways to choose k :
 - By cross-validation (or any validation)
 - Keep k odd to avoid equal voting for classes
 - Use \sqrt{n} where n is number of training instances